

Today's central concepts

• Markov Decision Process (MDP)

A Markov chain where the transition probabilities depend on the choice of an action. The goal is to find actions (a policy) that are optimal with respect to some objective function.

- State-space, S
- Actions
- Rewards (or costs)
- Transition probabilities: $P_t(s' | s, a)$
- Time-horizon and objective

- Finite-horizon, $T < \infty$

$$\mathbb{E} \left\{ \sum_{t=0}^{T-1} r_t(s_t, a_t) + r_T(s_T) \right\}$$

Need to define the terminal rewards, $r_T(s)$.

- Infinite horizon

- Discounted, $\mathbb{E} \left\{ \sum_{t=0}^{\infty} \gamma^t r_t(s_t, a_t) \right\}$

Two interpretations:

i, Value of unit reward decreases with time at geometric rate γ

ii, we are optimizing the total cost over a random time horizon. The system "shuts down" with probability $1-\gamma$ each time step.

- Average reward, $\lim_{T \rightarrow \infty} \mathbb{E} \left\{ \frac{1}{T} \sum_{t=0}^{T-1} r_t(s_t, a_t) \right\}$

Algorithms to solve MDPs:

(move next session end in computer labs)

- Dynamic programming, backward induction for finite horizon:

$$• u_T^*(s_T) = r_T(s_T)$$

$$• u_t^*(s_t) = \max_a \left\{ r_t(s_t, a) + \sum_{s' \in S} p_t(s' | s_t, a) u_{t+1}^*(s') \right\}$$

$$• a_t^*(s_t) \in \arg \max_a \left\{ \text{---} \parallel \text{---} \right\}$$

- Policy and value iteration
- Linear programming

Ex 2.1.1

b,

Try to pass x laws: l_1, l_2, \dots, l_x .

Every law doubles initial wealth w_0 .

Probability p_r of revolt \rightarrow lose everything.

Probability p_p of being rejected by parliament.

Can retire at any time. Want to maximize wealth. Model as MDP.

Solution:

We define the quantities that make up an MDP in order.

State-space:

Let

$$S = \{0, 1, \dots, N\} \cup \{\text{Retired}\} \cup \{\text{Fired}\}.$$

The set $\{0, 1, \dots, N\}$ of states denotes the number of laws that has been accepted.

The other two states represent if the dictator has retired or has been overthrown.

(Such states are called "auxiliary", "grave",
"terminal" or "absorbing" states.)

Actions:

He can do two things at every time:

R - retire , or T - try to pass
2 laws

Note:

The unit of time here is "number of times he has tried to pass a law", which is not necessarily equidistant in actual time.

reward/objective,

(There could be one year when two laws are up for vote, and some year when only one is.)

Time-horizon and cost function:

There is a total of N laws, so $T = N$.

The finite-horizon total reward objective is appropriate:

$$\mathbb{E} \left\{ \sum_{k=0}^{T-1} r_k(s_k, a_k) + r_T(s_T) \right\}.$$

Rewards:

Since there is a risk that he loses his full wealth, we can model it as if he collects the wealth only when he retires.

(Otherwise, we would have to define a cost, i.e., negative reward, that zeros what he has accumulated up to the point that a revolt happens.)

Terminal rewards:

- $r_T(s=n) = \omega_0 \cdot 2^n$
- $r_T(s=\text{Fired}) = 0$
- $r_T(s=\text{Retired}) = 0$


Rewards:

- $r_k(s = n, a = R) = \omega_0 \cdot 2^n$
- $r_k(s = \text{Fired}, a = \cdot) = 0$
- $r_k(s = \text{Retired}, a = \cdot) = 0$
- $r_k(s = n, a = T) = 0$

← collect reward equal to initial wealth doubled as many times as number of passed draws when he retires.

Transition probabilities:

It usually helps to draw the state-space to identify non-zero transition probabilities and only list those:



The non-zero transitions $p_t(s'|s, a)$ are:

- $P_t(s'=n+1 | s=n, a=\tau) = (1-p_r)(1-p_p)$

that the law
is passed successfully

tries
pass to
a law

no revolt
passed in
parliament

- $P_t(s'=n | s=n, a=\tau) = (1-p_r)p_p$

it did not
pass

- $P_t(s'=\text{Fired} | s=n, a=\tau) = p_r$

there was a revolt

Sanity check:

The sum of all outgoing transitions from any state should be one:

under
any
action

$$\sum_{s'} P_t(s'|s=n, a=\tau) = (1-p_r)(1-p_p) + (1-p_r)p_p + p_r =$$

$$= 1 - p_p - p_r + p_p p_r + p_p - p_r p_p + p_r =$$

$$= 1 \quad \text{oh!}$$

- $P_t(s' = \text{Retired} \mid s = n, a = R) = 1$

He decides to retire

- $P_t(s' = \text{Retired} \mid s = \text{Retired}, a = \cdot) = 1$

He stays retired

- $P_t(s' = \text{Fired} \mid s = \text{Fired}, a = \cdot) = 1$

He stays fired.

Ex 2.6 | The rational thief

He is caught with probability p , and loses everything. Otherwise, he collects the valuables and adds them to his fortune. He can retire at any time.

Solution:

State-space:

We cannot take # of houses robbed (as we did in previous exercise), since each house contains different valuables.

Let

$$S = R_{\geq 0} \cup \{ \text{Prison} \} \cup \{ \text{Retired} \}$$

where \nearrow is his accumulated fortune.

Actions:

Every night, he can do two things:

R - retire, or C - continue.

Rewards:

We let the thief actually receive his earnings when he chooses to retire:

- $r_k(s=x, a=R) = x$
- $r_k(s=x, a=C) = 0$
- $r_k(s=\text{Prison}, a=.) = 0$
- $r_k(s=\text{Retired}, a=.) = 0$

Intuition:
Due to the risk of him losing everything accumulated so far, his current wealth has to be a part of the system's state. At some point we will have to subtract this number.

Time-horizon and cost function:


There are multiple ways of modeling this problem.

- i) Essentially, as formulated, this is an infinite-horizon total reward problem:

$$\mathbb{E} \left\{ \sum_{k=0}^{\infty} r_k(s_k, a_k) \right\}.$$

However, this looks dangerous — will this sum be finite?

In our model, the system will always end up in a no-reward absorbing state:



So the terms in the sum are actually zero after some (random!) time.

Hence, in effect, this is a finite horizon problem (with random horizon). This criterion is valid, for example, if there exists a time such that for any initial state and policy, there is positive probability of ending up

in a reward-free absorbing state after this time.

(See Sec. 7.2 of "Dynamic Programming and Optimal Control" by Dimitri Bertsekas (vol. 1) for details.)

ii, we could assume that the thief is only fit to rob houses up to some age. we would then have a standard finite horizon problem:

$$\mathbb{E} \left\{ \sum_{t=0}^{T-1} r_t(s_t, a_t) + r_T(s_T) \right\}.$$

In this case, we need to define the terminal rewards:

- $r_T(x) = x$
- $r_T(\text{Prison}) = 0$
- $r_T(\text{Retired}) = 0$

← He requires what he has accumulated, if he has not yet retired

iii, One interpretation of the discounted infinite-horizon objective:

$$\mathbb{E} \left\{ \sum_{k=0}^{\infty} \gamma^k r_k(s_k, a_k) \right\}$$

is that we are actually optimizing a total reward (i.e., undiscounted) criterion, but that

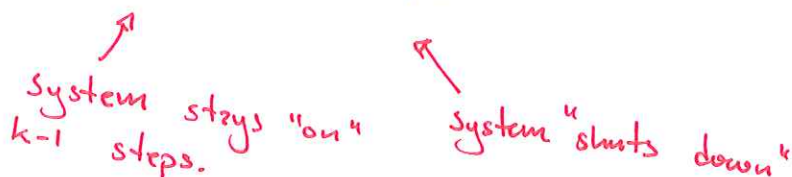
the system "shuts down" with probability $1-\lambda$ in each time step.

In other words, where we optimize

$$\mathbb{E} \left\{ \sum_{k=0}^T r_k(s_k, a_k) \right\}$$

for a T that is geometrically distributed:

$$\Pr\{T=k\} = \lambda^{k-1} \cdot (1-\lambda).$$


System stays "on" $k-1$ steps. System "shuts down"

The expected "life-time" of the system is

$$\mathbb{E}\{T\} = \frac{1}{1-\lambda}.$$


(This interpretation is equivalent to i, — see
Bertsekas' book Sec. 7.3 for details.)

Transition probabilities:

Assume we use i , or ii .

Let the value of the variables in the house robbed on night k be w_k .

It is convenient to draw the state-space and only write down the non-zero transitions:



- $P_t(s' = \text{Prison} \mid s = \text{Prison}, a = \cdot) = 1$
 - $P_t(s' = \text{Retired} \mid s = \text{Retired}, a = \cdot) = 1$
 - $P_t(s' = \text{Retired} \mid s = x, a = R) = 1$
 - $P_t(s' = \text{Prison} \mid s = x, a = C_1) = p$
 - $P_t(s' = y \mid s = x, a = C_1) = (1-p) \times P_r\{w_t = y-x\}$
- Annotations for the last equation:
- his new wealth \uparrow y
 - old wealth \uparrow x
 - he did not get caught \uparrow $(1-p)$
 - Probability that the house was worth $y-x$ \uparrow $P_r\{w_t = y-x\}$
- Annotations for other equations:
- He stays retired/in prison (for the first two equations)
 - He chose to retire (for the third equation)
 - He tried to rob a house but got caught (for the fourth equation)

Alternative solution using iii) -

We remove the prison state, since it is equivalent to the process ending.

State-space:

$$S' = R_{\geq 0} \cup \{\text{Retired}\}$$

Actions:

R - retire, C - continue

Rewards:

- $r_t(s=x, a=R) = x$

- all other zero:

$$r_t(s=x, a=C) = 0$$

$$r_t(s=\text{Retired}, a=\cdot) = 0$$

Time-horizon and cost function:

Discounted infinite horizon:

$$\mathbb{E} \left\{ \sum_{t=0}^{\infty} (1-p)^t r_t(s_t, a_t) \right\}$$


Transitions: (non-zero)

- $P_t(s'=\text{Retired} | s=\text{Retired}, a=\cdot) = 1$

- $P_t(s'=\text{Retired} | s=x, a=R) = 1$

- $P_t(s'=y | s=x, a=C) = \Pr\{\omega_t = y-x\}$

Interpretation:




with probability p of terminating
at each time step.

Ex 2.5 |

You are in a 100-story building and have 2 eggs. You want to determine, with the minimum number of drops in the worst case, from which floor it is safe to drop eggs.

Solution:

Notice the modular structure of the problem. If we drop an egg from floor a :



After each drop, we end up in the same situation as before, but with less floors.


This is typical of problems that can be solved using dynamic programming: the problem can be broken down into a collection of simpler subproblems. Each subproblem is solved only once and its solution stored.

The next time the same subproblem appears, instead of having to recompute its solution, we simply look it up!

Let $V(N, e)$ be the minimum number of drops needed if we have e eggs and N floors to test.

Then

$$V(N, e) = \min_{a=1, \dots, N} \left\{ 1 + \max \{ V(N-a, e), V(a-1, e-1) \} \right\}$$



if the egg does not break, we will use the minimum # of drops for the upper $N-a$ floors

if it breaks, we will require the minimum # of drops to check the lower $a-1$ floors with one egg less.

Since we consider the worst-case scenario, we take the maximum of these two.

Performing this test requires one drop, hence the $+1$.

Finally, we try to be optimal, so we minimize with respect to what floor a we drop it.

"Nature is allowed to choose floor and picks the one that always will require us to use as many drops as possible"


For the recursion to be well-defined, we need to define the base cases:

- $V(1, e) = 1$



if we are uncertain about one floor, we need to perform exactly one drop to be able to say if it is critical or not

- $V(N, 1) = N$



if we have only one egg, to be absolutely certain about which floor is critical, we need to start from the bottom and try all. The worst case is that the last is the critical one.

Evaluating the recursion with these base cases yields that

$$V(100, 2) = 14 \text{ drops.}$$

egg.py

```

### How NOT to implement it (solutions of subproblems are recomputed)

import numpy as np

N = 100
E = 2

def V(n, e):
    # Base cases
    if e == 1:
        return n

    if n == 1:
        return 1

    # The DP equation
    minimum = n+1;
    for a in range(1, n+1):
        minimum = min(minimum, max(V(n-a, e), V(a-1, e-1)) + 1)

    return minimum

print "Number of drops needed for %i floors and %i initial eggs:" % (N, E)
print V(N, E)

### How to implement it using dynamic programming (solutions of subproblems are
# stored and looked up).

import numpy as np

N = 100
E = 2

V_lookup = -np.ones((N, E))

def V(n, e):
    # Base cases
    if e == 1:
        return n

    if n == 1:
        return 1

    # Check if this value has already been computed
    if not V_lookup[n-1, e-1] == -1:
        return V_lookup[n-1, e-1]

    # The DP equation
    minimum = n+1; # Set to something high
    for a in range(1, n+1):
        minimum = min(minimum, max(V(n-a, e), V(a-1, e-1)) + 1)

    # Save the value
    V_lookup[n-1, e-1] = minimum

    return minimum

print "Number of drops needed for %i floors and %i initial eggs:" % (N, E)
print V(N, E)

### Output ###

# In [1]: %run egg.py
# Number of drops needed for 100 floors and 2 initial eggs:
# 14.0

```


Ex 2.2 | Have to sell apartment within N days. We receive an offer a_t every evening that has to be accepted or rejected the following day. The offers are multiples of 10 000 SEK, i.i.d., positive and upper-bounded. Once we sell, we get a daily interest rate $p > 0$ on the money we place in the bank.

Solution:

State-space:

The old offers are not relevant once we receive a new (since we cannot call back an old buyer). Hence, we let the state of the MDP be the current bid.

If $10\,000 \cdot x_{\max}$ is the highest possible bid, then

$$S = \{10\,000 \cdot x \text{ for } 0 \leq x \leq x_{\max}\} \cup \{\text{Sold}\}.$$

To simplify notation, let's agree that we always speak in units of 10k SEK, so that

$$S = \{0, 1, \dots, x_{\max}\} \cup \{\text{Sold}\}.$$

Actions:

Two choices:

A - Accept offer

R - Reject offer

Time-horizon and objective:

Finite-horizon N :

$$\mathbb{E} \left\{ \sum_{t=0}^{N-1} r_t(s_t, a_t) + r_N(s_N) \right\}$$

(x, x_{\max})
are
integers

Rewards:

Terminal:

- $r_u(\text{Sold}) = 0$

- $r_u(x) = x$

we are forced to sell at the end

Non-terminal:

- $r_t(s=x, a=A) = x \cdot (1+r)^{N-t}$

the current offer


future interest

- $r_t(s=\text{Sold}, a=A) = 0$

- $r_t(s=\text{Sold}, a=R) = 0$

- $r_t(s=x, a=R) = 0$

Transitions:



the next bid can be anything, so all these states are connected

The non-zero transitions are:

- $P_t(s'=x \mid s \neq \text{Sold}, a=R) = \Pr\{\omega_t = x\}$

if we reject, tomorrow's offer is drawn iid from the bids' distribution

- $P_t(s' = \text{Sold} \mid s \neq \text{Sold}, a = \text{sell}) = 1$

we decide to sell

- $P_t(s' = \text{Sold} \mid s = \text{Sold}, a = \bullet) = 1$

let's solve the MDP. Recall the backward induction:

- $u_N^*(s_N) = r_N(s_N)$
- $u_t^*(s_t) = \max_a \left\{ r_t(s_t, a) + \sum_{s' \in S} P_t(s' \mid s_t, a) u_{t+1}^*(s') \right\}$
- $a_t^*(s_t) = \arg \max_a \left\{ \text{---} \parallel \text{---} \right\}$

The base case:

$$u_N^*(\text{Sold}) = r_N(\text{Sold}) = 0$$

$$u_N^*(x) = r_N(x) = x$$

The recursion:

Consider first $s = \text{Sold}$:

$$u_t^*(\text{Sold}) = \max_a \left\{ \underbrace{r_t(\text{Sold}, a)}_{= 0 \text{ for any action}} + \sum_{s'} \underbrace{P_t(s' \mid \text{Sold}, a)}_{= 1 \text{ for } s' = \text{Sold}, \text{ zero otherwise}} u_{t+1}^*(s') \right\}$$

$$= \max_a \left\{ 0 + 1 \cdot u_{t+1}^*(\text{Sold}) \right\} =$$

$$= u_{t+1}^*(\text{Sold})$$

By induction, we conclude

$$u_t^*(\text{Sold}) = u_{t+1}^*(\text{Sold}) = \dots = u_N^*(\text{Sold}) = 0 \quad (*)$$

Consider $s = x$:

$$u_t^*(x) = \max_a \left\{ r_t(x, a) + \sum_{s'} p(s' | x, a) u_{t+1}^*(s') \right\} =$$

this is zero for
 $s' = \text{Sold}$ according
to $(*)$

$$= \max_a \left\{ r_t(x, a) + \sum_{y=0}^{x_{\max}} p(y | x, a) u_{t+1}^*(y) \right\} \quad (**)$$

let's evaluate the two actions separately:

If $a = A$:

$$\underbrace{r_t(x, A)}_{= x(1+r)^{N-t}} + \underbrace{\sum_{y=0}^{x_{\max}} p(y | x, A) u_{t+1}^*(y)}_{= 0 \text{ for all } y} = x(1+r)^{N-t}$$

"when we
accept, we
can only
go to
Sold"

If $a = R$:

$$\underbrace{r_t(x, R)}_{= 0} + \underbrace{\sum_{y=0}^{x_{\max}} p(y | x, R) u_{t+1}^*(y)}_{= \Pr\{\omega_t = y\}} =$$

$$\sum_{y=0}^{x_{\max}} \Pr\{\omega_t = y\} u_{t+1}^*(y) = \mathbb{E}_{\omega} \{ u_{t+1}^*(\omega) \}$$

Taken together into (**), we have that:

$$u_t^*(x) = \max \left\{ \underset{\text{accept}}{x(1+r)^{N-t}}, \underset{\text{reject}}{\mathbb{E}_\omega \{ u_{t+1}^*(\omega) \}} \right\}$$

$$= \max \left\{ \underset{\text{accept}}{x}, \frac{\underset{\text{reject}}{\mathbb{E}_\omega \{ u_{t+1}^*(\omega) \}}}{(1+r)^{N-t}} \right\}$$

$$\stackrel{\text{def.}}{=} \max \left\{ \underset{\alpha}{x}, \underset{\alpha}{\alpha_t} \right\}.$$

\Rightarrow (If we have sold, then the action we take is irrelevant.)

If at time t we receive offer x , then we should:

accept the offer if $x > \alpha_t$

reject the offer if $x \leq \alpha_t$

$$\text{for } \alpha_t = \frac{\mathbb{E}_\omega \{ u_{t+1}^*(\omega) \}}{(1+r)^{N-t}}.$$

if $x = \alpha_t$,
we can
do
whatever

This is a time-dependent threshold policy:

